

Pratkar — Business Requirements Document

Why this product should exist · Geek2Code 2026

2026-09-17 · @Someone

1. Executive Summary

India has a large, recurring transfer of money from households to insurers that neither party has tested, because testing it requires skills the household does not have.

In FY25, 87.50% of health insurance claims were settled by number, meaning roughly one in eight went unpaid (IRDAI, via Ministry of Finance, March 2026). Health insurance generated 61.5% of all Insurance Ombudsman complaints — 32,655 of 53,102. Of the complaints that reached a formal award, 50.8% were decided in the policyholder's favour. Very few rejections are contested; roughly half of those that are, succeed.

The constraint is not that policyholders lack rights. It is that exercising them requires reading a 40–80 page policy wording against a rejection letter and against IRDAI regulation. The only assisted route available today is paid human consultancy, which is bounded by expert hours and cannot reach the volume of rejections.

Pratkar converts that expert task into software: two documents in, a contestability verdict with a full evidence trail and a ready-to-file appeal out. It is free for individuals, and monetises the parties who handle rejections in volume — agents, hospital TPA desks, employer group-policy administrators and legal-aid organisations.

The competitive position is structural rather than technical. Every funded claims-AI company in India sells faster adjudication to hospitals and insurers, who are the counterparty to the user. A patient-side advocate is a business-model conflict for all of them.

This document defines why the product should exist. What it does is in the PRD; how it is built is in the Technology Stack and Architecture document.

2. Business Problem

Money that may be owed to households stays with insurers because the cost of checking exceeds what a household can pay.

The economics of the current situation

Party	Position today
Household	Bears the full loss of a rejected claim. Cannot assess the rejection. Cannot afford the assessment
Insurer	Faces challenge on a small fraction of rejections. Grounds that would not survive examination are rarely examined
Claim-advocacy firm	Has a viable business but a hard ceiling: revenue scales only with consultant headcount
Ombudsman	Receives a volume of complaints far below the volume of rejections, so its corrective function reaches few cases

Why the market has not corrected this. The value at stake in any single claim is too small to attract legal representation, and too large for the household to write off comfortably. That gap — individually modest, collectively very large — is precisely where software works and where per-case human labour does not.

The observable evidence. Health insurance produces 61.5% of all insurance grievances reaching the Ombudsman, the largest single category. Of complaints decided by formal award, 50.8% go in the complainant's favour. That success rate, against a complaint volume far below the rejection volume, indicates unexercised valid claims rather than a well-functioning correction mechanism. **[Deck]**

Cost to the household. Out-of-pocket expenditure remains 43.4% of India's total health spending (National Health Accounts 2022-23, released May 2026). A rejected claim converts insured expenditure into out-of-pocket expenditure, at the moment the household is least able to absorb it. **[Deck]**

3. Business Opportunity

The opportunity is to remove the expert-hour ceiling from claim advocacy.

Why solving this has value. The existing solution works — when a rejection is properly examined, roughly half are overturned. It simply cannot be delivered at the volume the problem occurs at, because each case consumes skilled human time. Software does not have that constraint. The same analysis that takes a consultant an hour can be performed repeatedly at near-zero marginal cost.

Who benefits

Beneficiary	Benefit
Policyholders	Recover claims they would otherwise absorb; avoid pointless appeals on valid rejections
Insurance agents and brokers	Answer client questions credibly without becoming claims experts
Hospital TPA and billing desks	Process rejections faster and with regulatory grounding
Employer group-policy administrators	Support employees at scale without an external consultancy contract
Legal-aid and consumer NGOs	Serve more complainants with the same staff
The wider market	Over time, evidence accumulates about which rejection grounds do not survive examination

What value is created. Three distinct kinds:

1. **Recovered value** — money returned to households on claims that should have been paid.
2. **Avoided waste** — time and cost saved when the system correctly says a rejection is valid and no appeal is worth filing. This is real value and is often overlooked.
3. **Capacity released** — expert hours freed at agencies, TPA desks and NGOs, redirected from routine assessment to the cases that genuinely need judgement.

Timing. Nothing in the product depends on a regulatory change, a partnership, or a platform launch. Policy wordings are published, IRDAI circulars are public, and the escalation process already exists. The opportunity is available now and does not require permission from any incumbent. **[Deck]**

4. Stakeholder Analysis

The user and the payer are different people, and that separation is deliberate.

Stakeholder	Category	Interest	Influence on the product
Policyholder	Primary user, not a payer	Recover a rejected claim; avoid a pointless appeal	Defines the entire primary flow

Stakeholder	Category	Interest	Influence on the product
Insurance agent or broker	Secondary user, prospective payer	Serve clients credibly; retain them	Drives Phase 2 volume features
Hospital TPA / billing desk	Secondary user, prospective payer	Resolve rejections at volume	Drives batch handling in later phases
Employer group-policy administrator	Secondary user, prospective payer	Support employees without a consultancy contract	Drives an organisational tier
Legal-aid / consumer NGO	Secondary user, prospective payer at concessional terms	Extend reach of limited staff	Drives accessibility and language coverage
Insurer / TPA	Counterparty , never a customer	Defend claim decisions	None. Selling to them would invert the product
IRDAI	Regulator whose published rules the product encodes	Policyholder protection	Source of the rulebook. Not approached, not integrated
Insurance Ombudsman	Institution the appeal escalates to	Resolve complaints	Process the output is formatted for. Not approached, not integrated
The build team	Decision maker for the hackathon	Ship a working, defensible MVP	Owens all scope decisions

Decision makers. For the hackathon, the five-member team alone. No external approval, licence or partnership is required to build or demonstrate the MVP. **[Deck]**

Paying customers. None in the MVP. The product is free for individuals, permanently, and revenue comes later from the volume handlers above. **[Deck]**

A stakeholder deliberately excluded. No insurer, TPA or claims-processing vendor can be a customer without compromising the user's interest. This is stated here so it is not revisited as an easy revenue idea later. **[Deck]**

5. Target Market

The addressable market has not been sized. What follows is what the source material supports, and what it does not.

Geography. India. Health insurance only for the MVP; motor and life are future scope.
[Deck]

What the source material establishes

Indicator	Figure	Source
Health claims settled by number, FY25	87.50%	IRDAI, via Ministry of Finance, March 2026
Health share of Insurance Ombudsman complaints, FY25	61.5% (32,655 of 53,102)	Council for Insurance Ombudsmen, Annual Report 2024-25
Complaints decided by formal award, FY25	18,171	Council for Insurance Ombudsmen, Annual Report 2024-25
Awards in the complainant's favour, FY25	9,234 (50.8%)	Council for Insurance Ombudsmen, Annual Report 2024-25
Out-of-pocket share of total health spending	43.4%	National Health Accounts 2022-23, released May 2026
Comparable service: cases handled since 2018	20,000+	Insurance Samadhan

Not established — requires research this project has not done

Unknown	Status
Total annual health claim rejections in absolute numbers	TBD
Rupee value of rejections that would be contestable	TBD
Number of insurance agents, TPA desks or employer administrators reachable as customers	TBD

Unknown	Status
Willingness to pay among any secondary user segment	TBD
Share of rejection letters that cite a clause in a machine-readable form	TBD — material to product viability, see PRD §4
Serviceable and obtainable market size	TBD

Position taken. This document does not estimate a market size. Deriving one from the figures above would require assumptions about rejection volumes and contestability rates that no source in our material supports. A sizing exercise is the first item of post-hackathon business work. **[TBD]**

6. Value Proposition

Customer problem	Product solution	Observable value
Cannot tell whether a cited clause applies	Retrieves the clause verbatim from their own policy and tests it against the policy and IRDAI provisions	A Strong, Moderate or Weak verdict with cited reasons, in minutes
Does not know that regulatory grounds exist that override policy wording	Runs coded IRDAI provisions before any interpretation	Invalid grounds surfaced that the policyholder could not have found
Cannot afford a consultant at the moment a claim is denied	Free for individuals	The assessment happens at all, rather than not happening
Does not trust an AI answer on a legal question	Every statement traces to a policy page or a named provision; the system abstains rather than guessing	A verdict the user can check themselves, line by line
Does not know how to write or file an appeal	Generates a grievance-officer letter from the established grounds	A document ready to send
Cannot read the policy or the process in English	Verdict and letter in Hindi	Access for users who could not otherwise use the product

Customer problem	Product solution	Observable value
Might waste effort appealing a valid rejection	Returns Weak verdicts and explains why the rejection stands	Time and effort saved on a case that would not succeed

The core proposition in one line. Pratkar converts a task that currently requires a paid expert into something a policyholder can complete in an evening, with every conclusion open to inspection.

What makes it defensible. Not the model — anyone can call a model. The defensible assets are the coded regulatory rulebook, the grounding discipline that makes the output checkable, and a market position no incumbent can occupy without conflicting with their existing customers. **[Deck]**

7. Business Model

Free for individuals, permanently. Revenue from the parties who handle rejections in volume. **[Deck]**

Why individuals never pay. A policyholder reaches this product at the point of a denied medical claim — often the worst financial moment of their year. Charging there would both suppress use and be wrong. It would also invite the obvious criticism from judges and mentors that the product monetises distress. **[Deck]**

Segment	Relationship	Proposed model	Status
Individual policyholders	Primary users	Free, permanently	Decided [Deck]
Insurance agents and brokers	Volume handlers	Per-seat subscription	Proposed [Rec]
Hospital TPA and billing desks	Volume handlers	Per-site licence	Proposed [Rec]
Employer group-policy administrators	Volume handlers	Organisational licence, banded by headcount	Proposed [Rec]
Legal-aid and consumer NGOs	Mission-aligned	Free or concessional	Proposed [Rec]

Pricing. No price point is set. The only reference in our material is the comparable human service, Insurance Samadhan, which charges ₹999 plus a 20% success fee per case. That is a benchmark for what the market already pays for this outcome, not a price we have

chosen. **[Deck]** Any figure quoted before customer conversations would be invention. **[TBD]**

Unit economics. Cost per analysis is dominated by model inference on a long policy document. Free-tier infrastructure covers the hackathon. Sustained free individual use requires either a subscription base that cross-subsidises it, grant or CSR funding, or both. Which of these carries the free tier is unresolved. **[TBD]**

Potential future revenue streams — all speculative, none validated:

1. An organisational tier with case management for volume handlers.
2. Anonymised, aggregated reporting on which rejection grounds do not survive examination — valuable to consumer bodies and researchers. Requires consent design and careful handling.
3. Motor and life claim rulebooks as additional product lines.

Revenue streams ruled out. Selling to insurers, TPAs or claims-processing vendors; commission on recovered amounts taken from the individual; advertising against users in financial distress. **[Rec]**

8. Business Requirements

ID	Business requirement	Rationale	Priority	Success criteria
BR-01	The product must be usable without payment by an individual policyholder	Charging at the point of a denied medical claim suppresses use and is ethically wrong	Must	No payment step exists anywhere in the individual flow
BR-02	Every conclusion must be checkable by the user against their own documents	Trust cannot be asserted in a legal context; it has to be demonstrable	Must	Every output statement resolves to a page span or a named provision
BR-03	The product must not act on the user's behalf in any legal process	Representation carries obligations a student team cannot meet	Must	The user files every document themselves; no automated submission exists
BR-04	The product must state when a rejection is valid	A tool that always recommends appealing is not advice and will lose credibility	Must	Weak verdicts are returned and explained, with no appeal generated

ID	Business requirement	Rationale	Priority	Success criteria
BR-05	No insurer, TPA or claims vendor may become a customer	Would invert the product's purpose and destroy its competitive position	Must	No revenue line targets a counterparty to the user
BR-06	The product must be usable by someone who does not read English	Language is a primary access barrier for the target user	Should	Verdict and appeal letter available in Hindi
BR-07	User documents must not be retained or used for training	Health and financial documents belonging to people in difficulty	Must	Documents discarded at end of session; provider training disabled
BR-08	Accuracy claims must be published with their sample size	An unqualified percentage from a small set misleads mentors, judges and users	Must	No percentage quoted without its denominator
BR-09	The regulatory rulebook must be maintainable by a non-developer	The rulebook is the defensible asset and its owner is not writing application code	Should	Rules are data, readable and extendable without touching application logic
BR-10	The product must degrade honestly when it cannot reach a conclusion	A confident wrong answer in this domain causes real harm	Must	"Not determinable" returned rather than a low-confidence verdict
BR-11	A route to revenue from volume handlers must remain open	Free individual use needs a sustaining source	Should	At least one secondary segment validated in conversation post-hackathon
BR-12	The product must be demonstrable end to end on real documents	Credibility with judges, mentors and early users depends on nothing being simulated	Must	Full flow runs live on documents not previously seen

9. Competitive and Alternative Landscape

Only the alternatives named in our source material appear below. No additional competitor research was performed for this document.

Alternative	What it is	Limitation
Do nothing — the default	The household absorbs the loss	The most common outcome, and the one the product competes against most often
Insurance Samadhan	India's leading claim-advocacy service; 20,000+ grievances since 2018 at ₹999 plus a 20% success fee [Deck]	Human consultants. Revenue scales only with headcount, so it cannot reach the volume of rejections
Self-filing	The policyholder writes their own grievance and escalates	Requires exactly the expertise the user lacks; this is the problem, not a solution
Engaging a lawyer	Formal representation	Cost usually exceeds the value of a single health claim
Funded claims-AI companies in India	Sell faster claim adjudication to hospitals and insurers [Deck]	Serve the counterparty. A patient-side advocate is a business-model conflict for them
General-purpose AI assistants	A user pastes their letter into a chatbot	No regulatory rulebook, no obligation to ground answers in the user's own policy, no citation the user can verify

The competitive position, stated plainly. Pratkar is not competing on model quality. It competes on three things the alternatives structurally cannot offer together: it is free at the point of need, it grounds every conclusion in documents the user can check, and it sits on the policyholder's side of the table. **[Deck]**

The honest vulnerability. Nothing prevents a well-resourced entrant from building the same thing. The barrier is not technical. It is that the organisations with the data, the distribution and the capital are the ones for whom this product conflicts with their existing customers. That is a real but not permanent advantage. **[Rec]**

10. Business Benefits

Benefits are stated as mechanisms. Where a figure would be required to quantify one, it is marked TBD rather than estimated.

Cost savings

- For the household: the cost of a paid consultant avoided — benchmark ₹999 plus a 20% success fee for the comparable human service. **[Deck]**
- For volume handlers: assessment time per case reduced. Magnitude **TBD** until measured with a real agent or TPA desk.

Efficiency improvement

- Analysis that takes a consultant an hour is performed in minutes, repeatedly, at near-zero marginal cost.
- Correctly identifying valid rejections prevents effort being spent on cases that would not succeed — an efficiency gain that only a system willing to say "no" can deliver.

Risk reduction

- For the household: the risk of missing the statutory deadline is reduced, because the deadline is surfaced with the verdict.
- For the household: the risk of filing a weak appeal and losing confidence in the process is reduced.
- For volume handlers: the risk of giving a client wrong guidance is reduced, because every statement is traceable to a source.

Revenue opportunities

- Subscription revenue from the segments in §7. All unvalidated. **TBD**
- For agents and brokers specifically, better claim-stage support plausibly improves client retention. Unmeasured. **TBD**

Operational benefits

- Expert hours at agencies, TPA desks and NGOs are redirected from routine assessment to cases needing judgement.
- The regulatory rulebook is maintained once and applies to every case, rather than living in individual consultants' experience.

User benefits

- A denied claim is assessed at all, rather than not being assessed.
- The assessment is checkable, so the user is not asked to trust an opaque answer on a consequential matter.

- Hindi output extends access to users who could not otherwise use the product.
[Deck]

Societal benefit

- Out-of-pocket health expenditure is 43.4% of India's total health spending. Every claim correctly recovered moves a household's cost back to the insurer it was insured with. Aggregate effect **TBD**. **[Deck]**

Sustainability note. The product has no material footprint beyond ordinary cloud compute. No sustainability claim is made, because none is supported.

11. Feasibility

Dimension	Assessment	Basis
Technical	High for the MVP. Document parsing, retrieval, a rule engine and constrained generation are all known-good components. No gated API, no model training, no simulated data	Every input is either supplied by the user or publicly published [Deck]
Operational	High for the hackathon, unproven beyond it. One session at a time, no accounts, no support burden. Running it for real users needs support, consent records and a person answering questions	MVP scope is deliberately single-session
Financial	High for the hackathon, unresolved beyond it. Free-tier infrastructure covers the build. Sustained free individual use needs a funding source that has not been chosen	See §7 unit economics [TBD]
Organisational	Adequate, with a known constraint. Two members build; three own the rulebook, design and testing. Those three workstreams are file-based and never block the code path	Team structure in the deck [Deck]

Dimension	Assessment	Basis
Scalability	Architecturally fine, commercially unproven. Each analysis is independent, so scaling is a cost question rather than a design question. Whether anyone will pay is untested	See §5 [TBD]
Regulatory	Workable as designed. The product prepares documents the user files themselves and gives no advice. It encodes published rules rather than interpreting law	Positioning decision, BR-03

The binding constraint is not technical. It is the supply of real rejection letters for testing and demonstration. No amount of engineering substitutes for them, and the MVP's credibility rests on running on genuine documents. **[Deck]**

The second constraint is commercial validation. Every revenue assumption in §7 is untested. The MVP can succeed as a hackathon project and still have no proven business, and this document does not claim otherwise. **[TBD]**

12. Business Risks

Technical and delivery risks are in PRD §18. These are the risks to the business case.

Category	Risk	Severity	Response
Market	No secondary segment proves willing to pay, leaving the free tier unfunded	High	Validate with real agents and TPA desks before building any paid tier. Grant and CSR funding as an alternative path [TBD]
Market	Rejection letters in practice are too vague to analyse, shrinking the addressable case base	High	Flow C handles it as a product feature. Measure the share of clause-citing letters in the labelled set — an early indicator of viability
Adoption	Policyholders never discover the product at the moment they need it	High	Distribution is unsolved. Agents, TPA desks and NGOs are as much a channel as a customer [TBD]

Category	Risk	Severity	Response
Adoption	Users do not trust an automated answer on a legal matter	Medium	The evidence trail is the response: the user checks the reasoning rather than trusting it
Regulatory	Positioning drifts from document preparation toward advice or representation	Medium	BR-03 fixes the boundary. No automated filing, no advice, user files their own documents
Regulatory	IRDAI provisions change and the rulebook goes stale	Medium	Rules are data with provision references attached, reviewable on a schedule. A stale rule is worse than no rule
Data	Insufficient real rejection letters for credible testing	High	The project's hardest dependency. Collect continuously; never make accuracy claims on constructed letters
Data	No lawful precedent corpus exists, capping verdict quality	Medium	Accepted and stated openly. Out of scope until a lawful source exists [Deck]
Privacy	A breach involving health and financial documents	High	Session-only retention, no training use, minimal logging. The strongest mitigation is holding almost nothing
Reputational	A wrong Weak verdict causes someone not to appeal, and they miss the deadline	High	Abstain rather than guess; always show the deadline; never state a verdict the evidence does not support
Execution	The team disperses after the hackathon and the project stops	Medium	Realistic. The documents and the coded rulebook are the durable assets either way
Competitive	A well-resourced entrant builds the same thing	Medium	The barrier is business-model conflict, not technology — real but not permanent [Rec]

13. Business KPIs

No targets are set for post-hackathon KPIs. The product has no users yet, so any number would be invented. **[TBD]**

Area	KPI	Why it matters	Measurable from
Adoption	Analyses completed	The base measure of whether anyone uses it	Launch
Adoption	Share of sessions reaching a verdict	Measures whether the two-document requirement is a real barrier	Launch
Usage	Appeal letters generated	Distinguishes browsing from acting	Launch
Usage	Share of verdicts that are Weak	A healthy system returns them; near-zero would indicate the product just says yes	Launch
Customer value	Appeals actually filed	The first real outcome measure	Requires follow-up with users
Customer value	Appeals succeeding	The outcome that matters most, and the hardest to observe	Requires follow-up; may never be fully measurable
Customer value	Amount recovered	The clearest expression of value created	Self-reported; treat as indicative
Operational	Grounding violations per thousand analyses	The integrity measure. Target is zero	Launch
Operational	Assessment time saved per case for volume handlers	The basis of any paid tier	Requires a pilot
Revenue	Paying organisations	Whether the model works at all	Post-validation
Retention	Volume handlers still active after 90 days	Whether it becomes part of their workflow	Post-launch

Area	KPI	Why it matters	Measurable from
Effectiveness	Verdict agreement with expert review	Whether the system is right, not just confident	Requires expert reviewer time

The one KPI with a target now. Grounding violations must be zero. Every other measure needs users, and the product does not have them yet.

Reporting discipline. Every figure published carries its sample size, per BR-08. A success rate from a handful of cases is reported as a count, never as a percentage.

14. Business Roadmap

Stage	Business objective	What must be true to move on
Hackathon prototype — by 20 Sep 2026	Prove the mechanism works on real documents and that the grounding discipline holds	PRD §17 acceptance criteria met
Pilot	Put it in front of a small number of real policyholders, with the team watching every case	Evidence that non-team users complete the flow and that verdicts survive expert review
MVP for real users	Operate continuously for individuals, free, with support and consent records in place	A measured accuracy figure the team will publish with its sample size; a decided answer on who funds the free tier
Production	Serve individuals reliably and validate a paid tier with one volume-handler segment	At least one organisation paying, and retention past 90 days
Commercial product	Sustain free individual use from subscription revenue; extend to motor and life	Unit economics that work without grant funding

Sequencing principle. Prove the product is right before making it bigger. Extending to new insurance lines multiplies the rulebook surface without strengthening the core; a second line before health is proven would weaken the thing that makes this defensible.

[Rec]

Decisions deferred by design. Pricing, the funding source for the free tier, distribution strategy, and whether to pursue any organisational tier at all. Each needs customer conversations the team has not had. Recording them as open is more useful than filling them with plausible numbers. **[TBD]**

Realistic note on continuity. The team is five students with a hackathon deadline. Whether this continues past 20 September is itself undecided. The coded regulatory rulebook and these documents are the assets that survive either way.